

# Exercise 3

Applications of Data Analysis

# 1 DATA PREPROCESSING

---

First of all the data is read with numpys genfromtxt – method into variables y and x. X contains modulator-data and y is metal concentration data.

```
basepath = os.path.dirname(__file__)
filepath = os.path.abspath(os.path.join(basepath, "Water_data.csv"))

y = np.genfromtxt(filepath, delimiter=',', skiprows=1, usecols=range(3, 6))
x = np.genfromtxt(filepath, delimiter=',', skiprows=1, usecols=range(0,3))
```

Data is standardized using z-score normalization.

```
def calculateZScore():
    xArr = np.asarray(x)
    zScores = (xArr - xArr.mean()) / xArr.std()
    return zScores
```

# 2 CALCULATING K NEAREST NEIGHBOR AND PREDICT VALUES

---

K nearest neighbor is calculated in inferErrors – method. This method calculates Euclidean distance for test instance against every instance of training set. Method returns k nearest neighbor leaving out n nearest neighbor (n = leaveOut attribute). In Leave-one-out the training set is same as test set and leaveOut is 1 to cut out the test instance from neighbors.

Values are predicted in chooseMajorityLabel – method. This method calculates mean for every modulators and returns the prediction based on that.

```

def chooseMajorityLabel(neighbors,k):
    predictedOutcome = []
    for i in range(3):
        sumOfMod = 0.0
        for j in range(len(neighbors)):
            sumOfMod = sumOfMod + neighbors[j][1][i]
        predictedOutcome.append(sumOfMod/k)

    return predictedOutcome

def inferNeighbors(trainSet,testInstance,labels,k,leaveOut):
    distances = []
    for x in range(len(trainSet)):
        distances.append((ssd.euclidean(trainSet[x], testInstance), labels[x]))

    distances.sort(key=operator.itemgetter(0))
    return distances[leaveOut:k+leaveOut]

```

### 3 LEAVE-**<T>**-OUT CROSS-VALIDATION

---

```

def LooCV(k):
    yPredictions = []
    stdX = calculateZScore()
    for i in range(len(stdX)):
        neighbors = inferNeighbors(stdX,stdX[i],y,k)
        yPredictions.append(chooseMajorityLabel(neighbors,k))

    cIndexCTotal = calculateCIndex(yPredictions,0,y)
    cIndexCd = calculateCIndex(yPredictions,1,y)
    cIndexPb = calculateCIndex(yPredictions,2,y)
    printCIndexes(cIndexCTotal,cIndexCd,cIndexPb)

```

LooCV (Leave-One-Out) calculates neighbors for every instance in stardardized training data and majority labels. Majority labels are included to predicted values of y.

The result using LOO:

Leave-one-out cross-validation  
C Indexes:  
c-total: 0.823497067449  
cd: 0.775997528249  
pb: 0.819797551789

```

def LfoCV(k):
    yPredictions = []
    stdX = calculateZScore()
    i = 0
    kf = KFold(len(stdX),67)
    for train_index, test_index in kf:
        xTrain, xTest = stdX[train_index], stdX[test_index]
        yTrain, yTest = y[train_index], y[test_index]
        for i in range(len(xTest)):
            neighbors = inferNeighbors(xTrain, xTest[i],yTrain, k)
            yPredictions.append(chooseMajorityLabel(neighbors, k))

    cIndexCTotal = calculateCIndex(yPredictions,0,yTrain)
    cIndexCd = calculateCIndex(yPredictions,1,yTrain)
    cIndexPb = calculateCIndex(yPredictions,2,yTrain)
    printCIndexes(cIndexCTotal,cIndexCd,cIndexPb)

```

LfoCV (Leace-Four-Out) divides data set into training and test sets. Test set contains for instances and every test set includes different instances. For every test set calculates neighbors and chooses majority label. Majority labels are included to predicted values of y.

The result using LFO:

```

Leave-four-out cross-validation
C Indexes:
c-total: 0.784681989924
cd: 0.664510447036
pb: 0.759032141131

```

## 4 CALCULATING C-INDEX

---

```
def calculateCIndex(predictions, index, labels):
    n = 0
    h_sum = 0
    for i in range(len(labels)):
        t = labels[i][index]
        p = predictions[i][index]
        for j in range(i+1, len(labels)):
            nt = labels[j][index]
            np = predictions[j][index]
            if t != nt:
                n = n + 1
                if (p < np and t < nt) or (p > np and t > nt):
                    h_sum = h_sum + 1
                elif (p < np and t > nt) or (p > np and t < nt):
                    h_sum = h_sum + 0
                elif (p == np):
                    h_sum = h_sum + 0.5

    if n == 0:
        return 0
    else:
        return h_sum/n
```

C-index values are calculated based on given pseudo-code.

## 5 CODE

---

'''

Authors: Marco Willgren, 502606  
Jarno Vuorenmaa, 503618

### Exercise 3: Applications of data analysis

The `Water_data.csv` file is a multi-parameter dataset consisting of 268 samples obtained from 67 mixtures of Cadmium, Lead, and tap water. Three features (attributes) were measured for each sample (`Mod1`, `Mod2`, `Mod3`).

#### Tasks

Use K-Nearest Neighbor Regression to predict total metal concentration (`c_total`), concentration of Cadmium (`Cd`) and concentration of Lead (`Pb`), for each sample.

- The data should be normalized using z-score.
- Implement Leave-One-Out Cross Validation approach and calculate the C-index for each output (`c-total`, `Cd`, `Pb`).
- Implement Leave-Four-Out Cross Validation and calculate the C-index for each output (`c-total`, `Cd`, `Pb`).

'''

```
import os
import numpy as np
import scipy.spatial.distance as ssd
import operator
from sklearn.cross_validation import KFold

if __name__ == '__main__':
    pass

basepath = os.path.dirname(__file__)
filepath = os.path.abspath(os.path.join(basepath, "Water_data.csv"))

y = np.genfromtxt(filepath, delimiter=',', skiprows=1, usecols=range(3, 6))
x = np.genfromtxt(filepath, delimiter=',', skiprows=1, usecols=range(0,3))

def LfoCV(k):
    yPredictions = []
    stdX = calculateZScore()
    i = 0
    kf = KFold(len(stdX),67)
    for train_index, test_index in kf:
        xTrain, xTest = stdX[train_index], stdX[test_index]
        yTrain, yTest = y[train_index], y[test_index]
        for i in range(len(xTest)):
            neighbors = inferNeighbors(xTrain, xTest[i],yTrain, k)
            yPredictions.append(chooseMajorityLabel(neighbors, k))

    cIndexCTotal = calculateCIndex(yPredictions,0,yTrain)
```

```

cIndexCd = calculateCIndex(yPredictions,1,yTrain)
cIndexPb = calculateCIndex(yPredictions,2,yTrain)
printCIndexes(cIndexCTotal,cIndexCd,cIndexPb)

def LooCV(k):
    yPredictions = []
    stdX = calculateZScore()
    for i in range(len(stdX)):
        neighbors = inferNeighbors(stdX,stdX[i],y,k)
        yPredictions.append(chooseMajorityLabel(neighbors,k))

    cIndexCTotal = calculateCIndex(yPredictions,0,y)
    cIndexCd = calculateCIndex(yPredictions,1,y)
    cIndexPb = calculateCIndex(yPredictions,2,y)
    printCIndexes(cIndexCTotal,cIndexCd,cIndexPb)

def printCIndexes(cTotal,cd,pb):
    print 'C Indexes:'
    print 'c-total: {a}'.format(a=cTotal)
    print 'cd: {b}'.format(b=cd)
    print 'pb: {c}'.format(c=pb)

def calculateCIndex(predictions,index,labels):
    n = 0
    h_sum = 0
    for i in range(len(labels)):
        t = labels[i][index]
        p = predictions[i][index]
        for j in range(i+1,len(labels)):
            nt = labels[j][index]
            np = predictions[j][index]
            if t != nt:
                n = n + 1
                if (p < np and t < nt) or (p > np and t > nt):
                    h_sum = h_sum + 1
                elif (p < np and t > nt) or (p > np and t < nt):
                    h_sum = h_sum + 0
                elif (p == np):
                    h_sum = h_sum + 0.5

    if n == 0:
        return 0
    else:
        return h_sum/n

def chooseMajorityLabel(neighbors,k):
    predictedOutcome = []
    for i in range(3):
        sumOfMod = 0.0
        for j in range(len(neighbors)):
            sumOfMod = sumOfMod + neighbors[j][1][i]
        predictedOutcome.append(sumOfMod/k)

    return predictedOutcome

```

```

def inferNeighbors(trainSet, testInstance, labels, k):
    distances = []
    for x in range(len(trainSet)):
        distances.append((ssd.euclidean(trainSet[x], testInstance), labels[x]))

    distances.sort(key=operator.itemgetter(0))
    return distances[1:k+1]

def calculateZScore():
    xArr = np.asarray(x)
    zScores = (xArr - xArr.mean()) / xArr.std()
    return zScores

LooCV(10)
LfoCV(10)

```